

Week 3 — Map It & Give It a Face

One-line claim, content map, identity kit, image curation

One-line claim

“I build AI systems the way a security audit would want them built.”

Chosen from 10 generated options, picked because it's a single, ownable stance rather than a feature list — it implies rigor without needing the word 'secure' to carry all the weight, and it sets up the 18/18 adversarial-test result as the proof rather than a decoration.

Content map

Single-page portfolio (one scrolling page, sectioned like a report) rather than a multi-page site — appropriate for one case study at this stage; each section below functions as the 'page' the brief asks for.

Section	Order	Content / case	Named CTA
Hero	1	One-line claim + credibility strip (18/18 tests passed)	Book a technical chat
Case study	2	Scholarship Agreement Assistant — strongest (only) case, leads immediately after hero	(reads on — no CTA, not yet earned)
Evaluation & Artifacts	3	Links to eval harness / repo so claims are checkable	View eval harness (repo link)
About	4	Technical point of view, not a bio	(reads on)
Contact	5	Single CTA repeated, email as secondary	Book a technical chat

All sections ladder to the same single action (book a technical chat) — no competing CTAs, per the earlier decision to make Contact single-CTA instead of email/scheduling split.

Still need to gather


- Real screenshots of the working prototype (login screen, chat view, Security Test tab with 18/18 result) — currently the page describes results in text/table form, not with real captures.
- Public repo URL for the eval harness / test suite, referenced in the Evaluation & Artifacts section but not yet linked to a live repo.
- Confirmed scheduling link (Calendly or equivalent) — CTA currently points to a placeholder anchor, not a working booking flow.

Identity kit

Type


One font: Inter, two weights (500 for body, 600–700 for headings/labels). No second display face — the earlier draft used three families (Space Grotesk, Inter, IBM Plex Mono); collapsed to one on purpose so nothing about the typography competes with the case study.

Palette

Role	Hex	Swatch
Background (near-black)	#14171C	
Text (near-white)	#EDEBE4	
Accent (single)	#C98A3E	

Note: this is a deliberate dark-mode inversion of the brief's literal 'near-black text on near-white background' — background and text are swapped, same principle (maximum, minimal contrast), chosen to read like an audit/verification log rather than a marketing page. Amber is the only accent; an earlier draft also used a semantic green for 'passed' test badges, removed so the amber flag/checkmark carries that meaning instead — one accent color, no exceptions.

Logo / favicon



A checkmark mark (amber on near-black) rather than initials alone — ties directly to the claim ('the way a security audit would want them built') without needing any illustration.

Style note

Fonts: Inter only (500 body, 600–700 headings/labels). **Colors:** background #14171C, text #EDEBE4, single accent #C98A3E — no second accent, no semantic colors. Mood: a security audit log, not a portfolio brochure — calm, flat, evidence-first.

Image curation

Images the content map actually needs:

- Login screen — real screenshot (own work, not generated)
- Chat view, scoped session — real screenshot (own work)
- Security Test tab showing 18/18 — real screenshot (own work, this is the single most important image on the page)
- Architecture diagram (login → scoped context → model → response) — technical diagram, already built as SVG; not decorative AI art, so it doesn't fall under the 'AI slop' judgment call the same way a generated hero image would
- Hero image — none. Deliberately zero decorative or generated imagery in the hero; a clean title over whitespace, per the brief's own 'reads as intentional' example
- About photo — optional, skipped for now (no real photo supplied yet); left as text-only rather than substituting a generated headshot-style image

Rejection note

No AI-generated image was produced or used anywhere on this page. The judgment call made here is the zero-image one: given the claim is about security rigor, any generated hero art (abstract circuits, glowing locks, glass-morphism shields — the standard 'AI security' visual cliches) would read as exactly the generic, unearned polish the whole page argues against. A real screenshot of a passing security test proves the claim; a generated padlock does not, no matter how polished it looks. That's the rejection: not one image out of a generated batch, but the entire category of image for this page.

Honesty note: no image-generation tool was available in this working session to produce an actual candidate image to reject side-by-side. If a real generated-and-rejected example is required for this deliverable, generate one or two hero options in an image tool and bring them back — the judgment above should hold, but it's worth verifying against a real image rather than this written argument alone.